

Bayesian Nonparametric Framework for Clustering and Generative Modeling of Multiple Time Series

Pierre le Roux*, Steve Kroon†

*Department of Computer Science, Stellenbosch University

Email: [REDACTED]

†Department of Computer Science, Stellenbosch University

Email: [REDACTED]

Abstract

This paper proposes a Bayesian nonparametric framework for clustering and generative modeling of multiple time series, integrating a Dirichlet Process Mixture Model (DPMM) for clustering with an Indian Buffet Process (IBP) for latent feature allocation within a state-space model. By avoiding predefined limits on the number of clusters or features, the approach can adapt its complexity to the data and discover shared temporal structures across series. Evaluation on synthetic datasets demonstrates that the model successfully captures broad clustering patterns and is capable of generating realistic synthetic time series, albeit with some limitations: threshold-based feature allocations often lead to oversimplified “all-on” or “all-off” usage, and certain parameters (e.g., process noise) are overestimated, reducing clustering accuracy (Adjusted Rand Index ≈ 0.35). Posterior predictive checks show moderate performance for most time series but underestimate high-variance behaviors. Future work includes refining prior specifications, exploring more robust inference methods, and testing on complex real-world data to fully exploit the model’s flexibility and improve its feature-allocation fidelity.

Index Terms

Bayesian Nonparametrics, Dirichlet Process, Indian Buffet Process, Time Series Clustering, Generative Modeling, Dirichlet Process Mixture Models

I Introduction

I-A Background and Motivation

In today’s data-driven world, analyzing and modeling multiple time series has become increasingly critical across diverse domains such as finance, healthcare, environmental monitoring, and social sciences. For example, financial analysts monitor numerous stock prices and economic indicators to identify patterns, predict trends, and make informed investment decisions [16]. In healthcare, tracking multiple patient vitals over time aids in diagnosing conditions and tailoring personalized treatments [10]. Effectively modeling and interpreting these multiple time series is essential for predictive analytics and strategic decision-making processes.

Traditional approaches to modeling multiple time series often employ mixture models, which assume that each time series belongs to one of a predefined number of clusters. While these models can capture inherent similarities among time series within the same cluster, they have significant limitations. Primarily, the necessity to specify the number of clusters *a priori* can lead to suboptimal modeling, especially when the true number of clusters is unknown or varies with the data [4]. Moreover, traditional mixture models lack the flexibility to dynamically accommodate the creation of new clusters as more data becomes available or as underlying patterns evolve over time, impeding the model’s ability to adapt to complex and changing data structures.

Bayesian nonparametric methods offer a promising alternative by providing models that can adapt their complexity to the data. Unlike parametric models, Bayesian nonparametrics do not require the specification of a fixed number of clusters or features beforehand. Instead, they allow for an infinite-dimensional parameter space, enabling the model to grow in complexity as needed [6]. The Dirichlet Process Mixture Model (DPMM) and the Indian Buffet Process (IBP) are two prominent Bayesian nonparametric methods facilitating flexible clustering and feature allocation, respectively. The stick-breaking construction of the DPMM allows for an unbounded number of clusters, dynamically introducing new clusters as required by the data [18]. Similarly, the IBP enables the allocation of an unbounded number of latent features to data points, capturing intricate dependencies and shared structures among multiple time series [8].

I-B Problem Statement

The primary objective of this research is to enhance existing generative models for individual time series by concurrently modeling multiple time series using Bayesian nonparametric methods. Specifically, this study aims to address the limitations of traditional mixture models in handling an unknown and potentially unbounded number of clusters and latent features. By leveraging the Dirichlet Process Mixture Model (DPMM) with a stick-breaking construction for cluster assignments and the

Indian Buffet Process (IBP) for latent feature allocation, we seek to develop a flexible and scalable framework capable of clustering multiple time series and inferring shared latent structures without prior knowledge of their quantities.

The scope of this study encompasses the development and implementation of the proposed Bayesian nonparametric model using synthetic datasets, facilitating controlled experimentation and validation. While the initial focus is on synthetic data generation and model validation, the framework is designed with scalability in mind, allowing for potential extensions to real-world datasets in future work. Limitations of this study include the computational complexity associated with Bayesian nonparametric methods, particularly in scenarios involving large-scale data, and the challenges inherent in accurately inferring latent structures from noisy and high-dimensional time series data.

I-C Contributions

This research offers several contributions to Bayesian nonparametric time series modeling and reproducible research practices:

- **Unified DPMM-IBP Framework for Time Series Clustering and Feature Allocation:** We propose a unified Bayesian nonparametric framework that combines the Dirichlet Process Mixture Model (DPMM) for clustering with the Indian Buffet Process (IBP) for feature allocation. This integration enables the simultaneous discovery of cluster structure and shared latent features across multiple time series, without requiring a priori limits on the number of clusters or features.
- **Robust Stick-Breaking Construction:** We develop and validate a stick-breaking approach for the DPMM that correctly handles leftover “stick” lengths, ensuring that mixture weights sum to one. This resolves the infinite log-probability issues arising from unallocated stick mass, thus facilitating stable, dynamic creation of new clusters.
- **Notebook-Based Synthetic Data Pipeline:** We implement a comprehensive pipeline (within a Jupyter or Colab notebook) for generating synthetic time series with known cluster assignments and latent feature structures. This design enables controlled experiments to assess the model’s capacity to uncover underlying patterns, while streamlining reproducibility and iterative experimentation.
- **Model Evaluation, Diagnostics, and Generative Checks:** We offer detailed convergence diagnostics (trace plots, $\hat{R}$, effective sample size), parameter recovery analyses, clustering performance metrics, and posterior predictive checks. These evaluations illustrate the generative capabilities of the model, reveal trade-offs in cluster/feature recovery, and provide guidelines for diagnosing inference quality.

I-D Organization of the Report

The report is structured as follows:

- **Introduction:** Introduces the research problem, objectives, and the motivation behind using Bayesian nonparametric methods for generative modeling of multiple time series.
- **Literature Review:** Provides an overview of existing methodologies related to time series clustering, Bayesian nonparametric methods, and generative models, highlighting relevant advancements and identifying gaps this study aims to address.
- **Methodology:** Details the synthetic data generation process and the construction of the Bayesian nonparametric model, elaborating on the theoretical underpinnings of the Dirichlet Process Mixture Model and the Indian Buffet Process.
- **Implementation:** Discusses the software tools and libraries utilized, the structure of the notebook-based workflow, and measures taken to ensure reproducibility, addressing challenges encountered during implementation.
- **Results and Evaluation:** Presents the outcomes of the analysis in several subsections:
 - **Synthetic Data Overview:** Describes the characteristics of the synthetic dataset and initial visualizations.
 - **Model Inference Parameters:** Examines the convergence and recovery of model parameters.
 - **Feature Allocation Results:** Evaluates how well the model allocates latent features.
 - **Clustering Results:** Assesses the clustering performance and alignment with true labels.
 - **Posterior Predictive Checks and Generative Capabilities:** Analyzes the model’s ability to replicate observed data and generate new series.
 - **Overall Summary of Results and Discussion:** Synthesizes key findings across parameters, features, clustering, and generative checks, interpreting the implications and trade-offs.
- **Limitations:** Identifies constraints and challenges faced during the study, including data complexity, inference difficulties, feature oversimplification, and computational burdens.
- **Future Work:** Outlines potential directions to refine the model, improve inference, and apply the framework to real-world datasets.
- **Conclusion:** Summarizes the key contributions, findings, and the significance of the research, highlighting areas for further development.
- **References:** Lists all cited works in IEEE format.
- **Appendices:** Provides supplementary materials such as code listings, detailed results, and additional figures.

This structured approach ensures a logical flow of information, enabling readers to grasp the foundational concepts, understand the methodological advancements, and appreciate the significance of the research findings within the broader context of Bayesian nonparametric modeling and time series analysis.

II Literature Review

II-A Time Series Clustering

Time series clustering involves grouping a set of time series such that series within the same cluster exhibit similar temporal patterns, while those in different clusters display distinct behaviors. This technique is pivotal in applications like financial market analysis, healthcare monitoring, and environmental studies [17]. Traditional clustering methods, such as K-means and hierarchical clustering, have been applied to time series data but often rely on predefined distance metrics and assume a fixed number of clusters, which may not capture the inherent temporal dependencies and dynamic nature of time series data [21].

One primary limitation of conventional clustering approaches is their inability to handle varying lengths and asynchronous time series. Additionally, these methods typically do not account for the temporal ordering of data points, potentially leading to suboptimal clustering outcomes [4]. To address these challenges, model-based clustering methods have been developed, assuming that data within each cluster are generated by a specific probabilistic model. Among these, state-space models have gained prominence due to their flexibility in modeling temporal dynamics [29].

State-space models provide a structured framework for representing the underlying processes that generate time series data. These models consist of the state equation, describing the evolution of hidden states over time, and the observation equation, linking the hidden states to the observed data [2]. In clustering, state-space models allow for the incorporation of temporal dependencies and can capture complex dynamics within each cluster. For instance, Fraley and Raftery [4] proposed a model-based clustering approach using Gaussian mixture models, where each cluster is characterized by its own state-space model, facilitating the identification of clusters based on both the mean and temporal covariance structure.

Despite their advantages, state-space model-based clustering methods still face challenges related to scalability and the necessity to specify the number of clusters in advance. These limitations have motivated the exploration of Bayesian nonparametric methods, offering more flexible and adaptive clustering mechanisms.

II-B Bayesian Nonparametric Methods

Bayesian nonparametric methods extend traditional Bayesian approaches by allowing models to grow in complexity with the data. Unlike parametric models, which require the specification of a fixed number of parameters or clusters, Bayesian nonparametric models can adapt their structure based on the observed data [6]. This flexibility makes them well-suited for applications where the true underlying structure is unknown or expected to be complex.

II-B1 Dirichlet Process Mixture Models

The Dirichlet Process Mixture Model (DPMM) extends finite mixture models, such as Gaussian Mixture Models (GMMs) and Finite Mixture of Multinomials, by allowing an infinite number of potential clusters, enabling the model to determine the appropriate number of clusters based on the data [18]. The stick-breaking construction is a fundamental component of the DPMM, providing a tractable way to represent the mixture weights.

In the stick-breaking process, mixture weights are generated by sequentially breaking a unit-length “stick” into pieces. Each piece corresponds to the weight of a cluster and is determined by Beta-distributed random variables. The weight for the k^{th} cluster is defined as:

$$\pi_k = \beta_k \prod_{j=1}^{k-1} (1 - \beta_j),$$

where $\beta_k \sim \text{Beta}(1, \alpha)$, and α is the concentration parameter controlling the number of clusters [18]. A higher α leads to a larger number of clusters, accommodating more complex data structures.

The DPMM has been applied to various domains, including image segmentation, natural language processing, and time series analysis [31]. Its ability to automatically infer the number of clusters makes it a powerful tool for exploratory data analysis and unsupervised learning tasks.

II-B2 Indian Buffet Process

While the DPMM is effective for clustering, it does not directly address feature allocation, where each data point can possess multiple latent features. The Indian Buffet Process (IBP) is a Bayesian nonparametric prior designed for binary latent feature models, enabling each data point to be associated with an arbitrary number of features [8].

The IBP is analogous to the Dirichlet Process in clustering but operates in feature space rather than partition space. In the IBP, the number of features is potentially infinite, and each data point can have a subset of these features. The generative

process of the IBP allows for the discovery of shared latent features across data points, facilitating the modeling of complex dependencies and interactions.

In time series modeling, the IBP can uncover underlying patterns or factors influencing multiple time series simultaneously. This shared feature structure enables the identification of common drivers or stimuli affecting different time series, enhancing interpretability and predictive power [32].

II-C Generative Models for Time Series

Generative models aim to learn the underlying distribution of data, enabling the synthesis of new data points resembling the original dataset. In time series, generative models must capture temporal dependencies and dynamic patterns to produce realistic synthetic sequences [7].

II-C1 Existing Approaches

Traditional methods include autoregressive models like ARIMA, which rely on linear assumptions and are limited in capturing complex nonlinear dynamics [17]. Recent advancements leverage deep learning architectures, such as Recurrent Neural Networks (RNNs) and Generative Adversarial Networks (GANs), to model intricate temporal dependencies [12], [27]. RNN-based models, particularly Long Short-Term Memory (LSTM) networks, have shown success in generating coherent time series data [12]. GANs offer a robust mechanism for generating high-fidelity time series by learning to produce data indistinguishable from real samples [22].

However, these deep learning-based methods often require large amounts of training data and are computationally intensive. They also lack the inherent probabilistic framework that facilitates uncertainty quantification and interpretability [20].

II-C2 Synthetic Data Generation

Synthetic data generation plays a crucial role in validating and benchmarking generative models, particularly when real-world data is scarce or sensitive. Generating realistic synthetic time series allows researchers to test model robustness, evaluate performance under controlled conditions, and explore the impact of varying data characteristics [24].

Frameworks like TSGM provide pipelines for generating synthetic time series with predefined characteristics, including cluster assignments and latent feature structures [30]. This facilitates systematic evaluation of generative models, ensuring they can accurately capture and reproduce underlying data dynamics.

In addition to model validation, synthetic data generation aids in understanding complex models, uncovering hidden patterns, and conducting sensitivity analyses to assess model stability [20].

II-D Integration of Bayesian Nonparametrics in Time Series Modeling

The integration of Bayesian nonparametric methods with generative models for time series clustering represents a growing area of research. Existing studies have explored various facets of this integration, contributing unique insights and methodologies.

II-D1 Comparative Analysis

Approaches like hierarchical Dirichlet processes (HDP) extend the DPMM to allow for shared clusters across data groups, enhancing flexibility in multi-group clustering scenarios [31]. Sticky HDP-HMMs introduce temporal dependencies into clustering, enabling the modeling of sequential data with persistent cluster assignments [3]. For feature allocation, extensions of the IBP accommodate temporal dependencies and dynamic feature generation; for example, the Temporal Indian Buffet Process (T-IBP) incorporates temporal dynamics into feature allocation, allowing features to evolve over time [32].

Comparatively, our approach uniquely combines the stick-breaking DPMM with the IBP within a state-space model framework for time series clustering. This integration facilitates simultaneous clustering and feature allocation, enabling the discovery of both group structures and shared latent features without predefined counts.

II-D2 Gaps in Literature

Despite advancements, several gaps remain. Existing models often treat clustering and feature allocation as separate processes, lacking a unified framework to infer both structures simultaneously. Additionally, applying these methods to generative modeling of multiple time series is limited, particularly in scenarios requiring the simultaneous discovery of dynamic clusters and shared latent features.

Moreover, the computational complexity of Bayesian nonparametric models poses challenges for scalability with large-scale time series data. Existing approaches may struggle to efficiently handle high-dimensional and long-duration time series, necessitating the development of more scalable inference algorithms.

This research aims to bridge these gaps by presenting a unified Bayesian nonparametric framework integrating DPMM and IBP within a state-space model for generative modeling of multiple time series. By leveraging synthetic data generation and implementing efficient inference techniques, we seek to enhance the flexibility, scalability, and interpretability of time series clustering models.

III Methodology

This section outlines the methodology employed in this research, encompassing the synthetic data generation process and the construction of the Bayesian nonparametric model. The methodology is structured into subsections addressing each critical component of the research framework.

III-A Synthetic Data Generation

A central objective of this research is to evaluate how effectively the proposed Bayesian nonparametric model identifies latent clusters and features when the ground truth is known. To this end, we generate a controlled synthetic dataset encompassing three main components: (1) cluster assignments, (2) feature allocations, and (3) a state-space process for time series simulation.

III-A1 Design Rationale

The synthetic data must satisfy several goals:

- 1) **Enforce Known Clusters:** Each time series should belong to one of a fixed number of clusters, enabling a direct comparison between true and inferred labels.
- 2) **Incorporate Latent Features:** Distinct clusters should exhibit different feature usage patterns - motivating the use of an Indian Buffet Process (IBP).
- 3) **Emulate Realistic Temporal Dynamics:** A state-space model ensures that each time series evolves under cluster- and feature-specific parameters, while still exhibiting noise typical of real-world data.

III-A2 Hyperparameters and Their Roles

Table I lists the key hyperparameters used throughout the data generation process. These values balance complexity and tractability while highlighting typical model behaviors.

TABLE I: Key Hyperparameters for Synthetic Data Generation

Parameter	Symbol/Name	Value(s)
Number of time series	N	50
Time steps per series	T	100
Number of true clusters	num_clusters	3
IBP truncation level	K_{trunc}	20
Cluster-specific IBP hyperpars	α_{clusters}	{1.0, 3.0, 5.0}
Std. of feature effects	τ_A	0.5
Noise for AR param.	σ_A	0.05
Process noise variance	Q_{true}	0.5
Observation noise variance	R_{true}	0.5
Observation coefficient	C_{true}	2.0

a) Motivation for the Hyperparameter Choices.

- A moderate dataset size ($N = 50$) and length ($T = 100$) balances clarity of cluster patterns with computational feasibility.
- Using different α_{clusters} (e.g., 1.0, 3.0, 5.0) across clusters encourages distinct feature activation probabilities.
- Process ($Q_{\text{true}} = 0.5$) and observation ($R_{\text{true}} = 0.5$) noise levels allow for moderate variability without making the time series purely deterministic or extremely noisy.
- The tanh transformation on autoregressive coefficients helps ensure stability ($|A_n| < 1$).

III-A3 Cluster Assignments and Feature Allocation

a) Cluster Assignments.

Each of the N time series is assigned to one of three true clusters (labelled 0, 1, 2), sampled uniformly or through any other desired distribution. Conceptually, we store a vector $\mathbf{c} = (c_1, c_2, \dots, c_N)$ indicating the cluster identity for each series.

b) Cluster-Specific Feature Probabilities.

To model feature usage diversity across clusters, we assume each cluster k draws feature probabilities from a Beta distribution:

$$p_{k,f} \sim \text{Beta}(\alpha_k, 1.0),$$

where $f \in \{1, \dots, K_{\text{trunc}}\}$ indexes the available features. A larger α_k produces higher average probabilities of activation, fostering the idea that different clusters might systematically prefer or ignore certain features.

c) IBP-Based Allocation.

For each time series n (belonging to cluster c_n) and each feature f , we sample

$$Z_{n,f} \sim \text{Bernoulli}(p_{c_n,f}),$$

forming a binary matrix $\mathbf{Z}_{\text{true}} \in \{0, 1\}^{N \times K_{\text{trunc}}}$. This matrix captures whether feature f is “on” or “off” for series n . Importantly, the presence of cluster-specific probabilities ensures that each cluster exhibits a distinct pattern of feature usage.

III-A4 Sampling Feature Effects and AR Coefficients

a) Cluster-Specific Feature Effects.

We next define a matrix of feature effects $\lambda_A \in \mathbb{R}^{K \times K_{\text{trunc}}}$, where each row $\lambda_{A,k}$ corresponds to cluster k . The elements of each row are drawn from a normal distribution:

$$\lambda_{A,k,f} \sim \mathcal{N}(0, \tau_A^2),$$

indicating how feature f (when active) modifies the latent state evolution for cluster k .

b) Autoregressive Coefficients.

For each time series n , we collect the feature effects relevant to its cluster c_n . Specifically,

$$\theta_n = \sum_{f=1}^{K_{\text{trunc}}} Z_{n,f} \lambda_{A,c_n,f} + \epsilon_{A,n},$$

where $\epsilon_{A,n} \sim \mathcal{N}(0, \sigma_A^2)$ introduces slight variability in the sum. To ensure stationarity, we map θ_n via

$$A_n = \tanh(\theta_n),$$

so that $|A_n| < 1$ in most cases, reducing the risk of explosive growth.

III-A5 State-Space Simulation

We simulate each time series under the following state-space model:

a) State Equation.

$$x_{n,t} = A_n x_{n,t-1} + \epsilon_{n,t}, \quad \epsilon_{n,t} \sim \mathcal{N}(0, Q_{\text{true}}). \quad (1)$$

We initialize $x_{n,0} \sim \mathcal{N}(0, 1)$, giving each series a random starting state. The process noise variance Q_{true} determines how much each next state deviates from the autoregressive prediction.

b) Observation Equation.

$$y_{n,t} = C_{\text{true}} x_{n,t} + \nu_{n,t}, \quad \nu_{n,t} \sim \mathcal{N}(0, R_{\text{true}}). \quad (2)$$

An observation coefficient $C_{\text{true}} = 2.0$ scales the hidden state, and the observation noise variance $R_{\text{true}} = 0.5$ introduces moderate measurement uncertainty.

III-A6 Visualizing the Synthetic Dataset

a) Feature Allocation Heatmap.

To verify that clusters differ in their feature usage, we generate a heatmap of $\mathbf{Z}_{\text{true}}$. Each row represents one time series, annotated with its cluster label, while columns correspond to features. Distinct patterns of “on/off” features become evident across clusters.

b) Distribution of A_n .

We also plot a histogram of $\{A_n\}_{n=1}^N$ to confirm that the $\tanh$ transformation keeps the coefficients within a stable range. This helps ensure each time series remains bounded, reflecting typical autoregressive behavior.

c) Sample Time Series.

Finally, we visualize a few of the 50 generated series to gain intuition about their trajectories. Series from different clusters often display characteristic amplitudes or decay rates, consistent with the cluster-specific A_n and the chosen noise levels.

III-A7 Summary of Data Generation Choices

Overall, the synthetic data design combines:

- **Cluster-based Feature Probabilities** (modulated by α_{clusters}),
- **Additive Feature Effects** (mapped through $\tanh$) to produce stable autoregressive coefficients, and
- **Moderate Process/Observation Noise Levels** ($Q_{\text{true}} = R_{\text{true}} = 0.5$)

to create a testbed that challenges the model to recover both cluster boundaries and latent features. Later sections compare the inferred results - cluster assignments, feature allocations, and posterior predictive distributions - to these known ground truths, thereby gauging the effectiveness of the proposed framework.

III-B Bayesian Nonparametric Modeling

The core of our methodology lies in the integration of Bayesian nonparametric techniques to facilitate flexible clustering and feature allocation within a state-space model framework. Specifically, we employ a **Dirichlet Process Mixture Model (DPMM)** [18], [23] for clustering and an **Indian Buffet Process (IBP)** [8] for latent feature allocation. These components are seamlessly embedded within a state-space model [2] to capture both global cluster structures and local, feature-driven temporal dynamics of multiple time series.

III-B1 Dirichlet Process Mixture Model (DPMM)

The DPMM serves as the backbone for our clustering mechanism, enabling the model to infer an unknown number of clusters from the data without imposing a fixed upper bound. By allowing the data to guide how many clusters appear, the DPMM is well suited for scenarios where the true cluster count is initially unknown [23].

a) Stick-Breaking Construction

To realize the DPMM, we utilize the **stick-breaking process** [18], a constructive representation that generates an infinite sequence of mixture weights $\{\pi_k\}_{k=1}^{\infty}$ for the clusters. In practice, we truncate this infinite process to a finite level $K_{\max}$ for computational feasibility.

The stick-breaking process is defined as follows:

$$\beta_k \sim \text{Beta}(1, \alpha_{\text{dp}}), \quad k = 1, 2, \dots, K_{\max}, \quad (3)$$

$$\pi_k = \beta_k \prod_{j=1}^{k-1} (1 - \beta_j), \quad k = 1, 2, \dots, K_{\max}-1, \quad (4)$$

$$\pi_{K_{\max}} = \beta_{K_{\max}} + \prod_{j=1}^{K_{\max}} (1 - \beta_j). \quad (5)$$

Here, α_{dp} is the concentration parameter of the Dirichlet Process, controlling the likelihood of creating new clusters. A higher α_{dp} favors the generation of more clusters. The truncation at $K_{\max}$ ensures that the sum of the mixture weights $\sum_{k=1}^{K_{\max}} \pi_k = 1$, accommodating any residual stick by absorbing it into the final cluster's weight $\pi_{K_{\max}}$.

b) Cluster Assignments

Each time series $n = 1, 2, \dots, N$ is assigned to one of the $K_{\max}$ clusters based on the mixture weights $\boldsymbol{\pi} = (\pi_1, \pi_2, \dots, \pi_{K_{\max}})$:

$$c_n \sim \text{Categorical}(\boldsymbol{\pi}), \quad n = 1, 2, \dots, N. \quad (6)$$

The cluster assignment $c_n \in \{1, 2, \dots, K_{\max}\}$ determines the latent characteristics and feature allocations specific to each time series.

c) Implementation Considerations

In our Bayesian framework, implemented using PyMC , the concentration parameter α_{dp} is treated as a random variable with a prior distribution:

$$\alpha_{\text{dp}} \sim \text{Gamma}(1.0, 1.0). \quad (7)$$

This non-informative prior allows the data to primarily influence the inferred number of clusters. The stick-breaking weights β_k are sampled from their Beta distributions as defined in Equation (3). The mixture weights π_k are then deterministically derived from the sampled β_k values using Equations (4) and (5). Finally, each time series is assigned to a cluster c_n based on the categorical distribution parameterized by $\boldsymbol{\pi}$.

III-B2 Indian Buffet Process (IBP) for Feature Allocation

To model the latent feature structures within each cluster, we employ the **Indian Buffet Process (IBP)** [8]. The IBP is well-suited for scenarios where the number of features is potentially unbounded, allowing the model to infer the appropriate number of latent features from the data.

a) IBP Generative Process

The IBP defines a prior over the binary feature allocation matrix $\mathbf{Z} \in \{0, 1\}^{N \times F}$, where F is the truncation level representing the maximum number of possible features. Each entry Z_{nf} indicates the presence (1) or absence (0) of feature f in time series n .

The generative process for $\mathbf{Z}$ is as follows:

$$v_{kf} \sim \text{Beta}\left(\frac{\alpha_p}{F}, 1\right), \quad k = 1, 2, \dots, K_{\max}; \quad f = 1, 2, \dots, F, \quad (8)$$

$$Z_{nf} \sim \text{Bernoulli}(v_{c_n f}), \quad n = 1, 2, \dots, N; \quad f = 1, 2, \dots, F. \quad (9)$$

Here, α_p is the IBP hyperparameter influencing the expected number of features per cluster. The cluster-specific feature probabilities v_{kf} are sampled from Beta distributions, allowing different clusters to have varying propensities for activating features. Subsequently, each feature allocation Z_{nf} is determined by a Bernoulli trial with probability $v_{c_n f}$, ensuring that feature usage is both cluster-dependent and stochastic.

b) Implementation Considerations

Within the PYMC framework, the IBP is implemented by:

- 1) **Sampling Feature Probabilities:** Each cluster k and feature f has an associated feature probability v_{kf} , drawn from a Beta distribution with shape parameters $\frac{\alpha_p}{F}$ and 1. This formulation ensures that the probability of feature activation is influenced by the IBP hyperparameter α_p , which is itself modeled with a prior:

$$\alpha_p \sim \text{Gamma}(2.0, 1.0). \quad (10)$$

- 2) **Generating Feature Allocations:** For each time series n and feature f , the binary allocation Z_{nf} is sampled from a Bernoulli distribution parameterized by the cluster-specific probability $v_{c_n f}$:

$$Z_{nf} \sim \text{Bernoulli}(v_{c_n f}), \quad n = 1, 2, \dots, N; \quad f = 1, 2, \dots, F. \quad (11)$$

This process results in a feature allocation matrix $\mathbf{Z}$ that should reflect both the inherent feature diversity within each cluster and the overall sparsity dictated by the IBP prior.

III-B3 Integration with the State-Space Model

The DPMM and IBP components are integrated within a state-space model to capture the temporal dynamics of each time series influenced by its cluster-specific features [2].

a) Cluster-Specific Feature Effects

Each cluster k is associated with a vector of feature effects $\lambda_k \in \mathbb{R}^F$. These effects modulate the autoregressive dynamics of the time series belonging to cluster k . The feature effects are sampled as:

$$\lambda_k \sim \mathcal{N}(\mathbf{0}, \tau_A^2 \mathbf{I}), \quad k = 1, 2, \dots, K_{\max}, \quad (12)$$

where τ_A is the standard deviation hyperparameter controlling the variability of feature effects across clusters.

b) Autoregressive Coefficients

For each time series n , the autoregressive coefficient A_n is determined by the cumulative influence of its active features and a small stochastic perturbation:

$$A_n = \tanh\left(\sum_{f=1}^F Z_{nf} \lambda_{c_n, f} + \eta_n\right), \quad \eta_n \sim \mathcal{N}(0, \sigma_A^2). \quad (13)$$

The tanh function ensures that A_n remains within $(-1, 1)$, promoting stable and stationary time series dynamics. The additive noise term η_n introduces slight variability in the autoregressive coefficients, reflecting real-world uncertainties.

III-B4 State-Space Model for Inference

The temporal evolution of each time series is modeled using a state-space framework comprising a state (transition) equation and an observation equation [2]. This structure allows the model to capture both the hidden states driving the time series and the observable measurements influenced by these states.

a) Model Specification

State Equation:

$$x_{n,t} = A_n x_{n,t-1} + \epsilon_{n,t}, \quad \epsilon_{n,t} \sim \mathcal{N}(0, Q), \quad (14)$$

Here, $x_{n,t}$ represents the hidden state of time series n at time t , and A_n is the autoregressive coefficient specific to n . The process noise $\epsilon_{n,t}$ introduces variability into the state transitions, with variance Q controlling the magnitude of this noise.

Observation Equation:

$$y_{n,t} = C x_{n,t} + \nu_{n,t}, \quad \nu_{n,t} \sim \mathcal{N}(0, R). \quad (15)$$

In this equation, $y_{n,t}$ is the observed value of time series n at time t , C is the global observation coefficient, and $\nu_{n,t}$ represents the observation noise with variance R . The coefficient C scales the influence of the hidden state on the observed data, while R controls the measurement uncertainty.

b) Implementation Details

Initial States:

The initial hidden state $x_{n,0}$ for each time series n is sampled from a standard normal distribution:

$$x_{n,0} \sim \mathcal{N}(0, 1), \quad n = 1, 2, \dots, N. \quad (16)$$

This provides a random starting point for each time series, ensuring diversity in their temporal evolution.

Noise Parameters:

The process noise variance Q and observation noise variance R are modeled with Half-Normal priors to enforce non-negativity:

$$Q \sim \text{HalfNormal}(1.0), \quad (17)$$

$$R \sim \text{HalfNormal}(1.0). \quad (18)$$

The choice of Half-Normal distributions reflects a belief in moderate noise levels, preventing overly large variances that could obscure underlying signal patterns.

State-Space Dynamics Implementation:

Within PyMC, the state-space model is implemented using the `scan` function, which iteratively applies the state equation to compute hidden states $x_{n,t}$ for each time series n over all time steps $t = 1, 2, \dots, T$. This recursive computation efficiently handles the temporal dependencies inherent in the state-space structure.

c) Rationale for Model Choices

Several key decisions underpin the structure and components of our Bayesian nonparametric state-space model:

- **Nonparametric Clustering with DPMM:** The Dirichlet Process Mixture Model [18], [23] offers flexibility in discovering the number of clusters inherent in the data, a crucial advantage when prior knowledge is limited or when the true number of clusters is expected to vary.
- **Latent Feature Allocation with IBP:** The Indian Buffet Process [8] facilitates the identification of shared latent features across time series without predefining their number. This adaptability is essential for modeling complex, multifaceted temporal dependencies.
- **State-Space Framework:** Incorporating a state-space model [2] allows the capture of temporal dynamics through hidden states, providing a robust mechanism for modeling both the underlying processes driving the time series and the observable manifestations influenced by these processes.
- **Stability through tanh Transformation:** The application of the tanh function to the autoregressive coefficients A_n ensures that the hidden states remain bounded, promoting stationarity and preventing runaway dynamics in the time series.
- **Hierarchical Priors:** Modeling hyperparameters (α_{dp} and α_p) with Gamma priors introduces hierarchical structure, allowing the model to adapt its complexity based on the data rather than relying on fixed hyperparameter values.

III-B5 Inference via Markov Chain Monte Carlo (MCMC)

To perform posterior inference over the model parameters, we employ **Markov Chain Monte Carlo (MCMC)** sampling, specifically the **No-U-Turn Sampler (NUTS)** [13], an adaptive variant of Hamiltonian Monte Carlo (HMC). Rather than maximizing a single objective function, PyMC samples from the *joint posterior* of all latent variables, reflecting the Bayesian paradigm where:

$$p(\theta \mid \text{data}) \propto p(\text{data} \mid \theta) p(\theta),$$

where θ includes all parameters: $\{\alpha_{dp}, \beta_k, c_n, \alpha_p, v_{kf}, Z_{nf}, \lambda_k, x_{n,t}, Q, R, \dots\}$. PyMC's sampling aims to approximate this posterior distribution rather than optimize it. This approach provides full uncertainty quantification, since each θ -sample is drawn in proportion to the posterior density.

a) Sampling Configuration

Our MCMC setup is configured with the following parameters to balance computational efficiency and sampling accuracy:

- **Number of Draws:** 1000
We collect 1000 posterior samples to ensure a large enough sample size for accurate posterior summaries and diagnostics.
- **Number of Tuning Steps:** 1000
An equal number of tuning steps allows the sampler to adapt its step size and mass matrix, enhancing convergence and reducing autocorrelation in the samples.

- **Number of Chains:** 2

Running multiple chains facilitates convergence diagnostics and ensures that the sampler explores different modes of the posterior.

- **Target Acceptance Rate:** 0.95

A higher target acceptance rate reduces the likelihood of rejection during sampling, promoting efficient exploration of the posterior [13].

b) Implementation Process

The inference process is executed within the PyMC framework as follows:

- 1) **Model Specification:** All model components - including DPMM, IBP, and state-space dynamics - are defined within a single PyMC model context.
- 2) **Sampling:** The `pm.sample` function is invoked with the specified configuration parameters, initiating the NUTS sampler to draw samples from the posterior distribution.
- 3) **Trace Extraction:** The resulting trace, stored in `trace_joint_model`, contains samples for all stochastic variables, including cluster assignments (c_n), feature allocations (Z_{nf}), autoregressive coefficients (A_n), and noise parameters (Q and R), among others.

c) Diagnostics and Convergence Assessment

To ensure the reliability and validity of the posterior samples, we perform several diagnostic checks:

- **Trace Plots:** Visual inspection of trace plots for each parameter assesses whether chains have adequately mixed and converged. Well-mixed traces exhibit “hairy caterpillar” appearances without evident drifts.
- **Potential Scale Reduction Factor ($\hat{R}$):** The $\hat{R}$ statistic measures the ratio of between-chain to within-chain variance. Values near 1 indicate convergence [5]. Typically, $\hat{R} < 1.1$ is considered acceptable.
- **Effective Sample Size (ESS):** ESS quantifies how many independent samples the chain effectively yields. Higher ESS values imply better sampling efficiency and more reliable posterior estimates.
- **Posterior Predictive Checks (PPC):** Utilizing the `ArviZ` library [19], we perform PPC to compare observed data with samples from the posterior predictive distribution, validating model fit.
- **Autocorrelation Analysis:** Examining autocorrelation helps detect persistent dependencies, indicating inadequate mixing or multiple posterior modes.

d) Tools and Visualization

We leverage the following tools for diagnostics and visualization:

- **ArviZ Library [19]:** Provides diagnostic summaries, trace plots, and posterior predictive checks for PyMC traces.
- **Matplotlib:** Used for custom plots (e.g., histograms, trace plots) to supplement `ArviZ` outputs.
- **Chain Splitting and Parallelization:** Running multiple chains in parallel accelerates sampling and ensures broader posterior exploration.

e) Rationale for MCMC Choices

NUTS (the No-U-Turn Sampler) [13] can handle complex, high-dimensional posteriors without extensive manual tuning of path lengths. Its adaptive nature reduces autocorrelation and increases sampling efficiency. Configuring a high target acceptance rate (0.95) further lowers rejection rates and yields smoother posterior exploration.

f) Integration of Model Components

The seamless integration of DPMM, IBP, and the state-space model within a single Bayesian framework allows for joint inference over cluster assignments, feature allocations, and temporal dynamics. This holistic approach ensures that dependencies between clusters, features, and time series evolution are coherently captured, enabling the model to leverage shared structures and latent patterns effectively.

g) Summary of Bayesian Nonparametric Modeling

In summary, our Bayesian nonparametric modeling approach leverages the flexibility of DPMM and IBP [8], [18], [23] to uncover clusters and latent features without fixed bounds, while the state-space model [2] captures temporal dependencies. Under PyMC’s paradigm, we define a joint likelihood over all unknowns and sample from the posterior $\prod p(\theta)p(\text{data} \mid \theta)$, rather than maximizing it. Through MCMC (NUTS) [13], we obtain a posterior distribution that adapts to the data’s complexity and provides a rich, uncertainty-aware representation of clustering, feature usage, and dynamic states in multiple time series.

IV Implementation

This section details the practical considerations and code organization underlying our Bayesian nonparametric state-space model, complementing the methodological framework described in Section III. It covers the computational environment, software tools, overall code structure, and reproducibility measures.

IV-A Software and Computational Environment

All implementations were carried out in **Python** (Version 3.10.8), a versatile language widely used for machine learning and data science [33]. The following core libraries facilitated various stages of model development, inference, and visualization:

- **PyMC** (Version 4.0.0) [28]: For specifying and performing Bayesian inference on our hierarchical models, leveraging the No-U-Turn Sampler (NUTS).
- **NumPy** (Version 1.23.5) [9]: For efficient numerical operations and array manipulations.
- **Matplotlib** (Version 3.6.2) [15]: For generating plots of synthetic data, diagnostics, and result visualizations.
- **ArviZ** (Version 0.12.1) [19]: For posterior analysis, trace plots, and convergence diagnostics.
- **PyTensor** (Version 2.8.11): Providing the computational graph layer underpinning PyMC’s symbolic operations.

All libraries were installed via `pip` to ensure straightforward reproducibility. Model training, sampling, and evaluations ran on an Intel Core i7 processor with 16 GB RAM, typically on Google Colab’s CPU runtime.

IV-B Code Organization in a Notebook Environment

Rather than developing a traditional multi-module Python codebase, we adopted a Jupyter (or Google Colab) notebook workflow to integrate data generation, model specification, inference, and evaluation in a single interactive document. This approach facilitates rapid experimentation and clear traceability of each computational step. The notebook is logically segmented into the following parts:

- **Data Generation Cells:** Define all synthetic data parameters, cluster assignments, and feature allocations (IBP), then simulate time series under the state-space model.
- **Model Specification Cells:** Construct the Bayesian nonparametric state-space model using PyMC, incorporating the stick-breaking DPMM, IBP feature allocations, and the state/observation equations.
- **Inference Cells:** Run Markov Chain Monte Carlo (MCMC) with the No-U-Turn Sampler (NUTS), specifying the number of draws, tuning steps, and acceptance rate. Posterior samples are stored in a trace object for subsequent analysis.
- **Evaluation and Visualization Cells:** Perform convergence diagnostics using ArviZ, generate posterior predictive checks, plot cluster assignments, and visualize time series reconstructions.

This notebook-based structure streamlines the development cycle, allowing immediate inspection of intermediate results and adjustments to hyperparameters or model components. To reproduce our work, readers can run the notebook in Google Colab using the provided environment specifications.

IV-C Reproducibility Measures

Scientific reproducibility was a principal concern throughout implementation:

IV-C1 Random Seed Initialization

A fixed random seed (`np.random.seed(42)`) was used to initialize all stochastic components, including data generation and MCMC initialization. This ensures that repeated runs yield identical synthetic datasets and starting points, facilitating consistent comparisons.

IV-C2 Code and Data Availability

All code, contained within Jupyter notebooks, along with synthetic datasets, is accessible in a public GitHub repository.¹ The repository includes:

- `main_notebook.ipynb`: The primary Jupyter notebook integrating data generation, model specification, inference, and evaluation.
- `README.md`: Step-by-step instructions for environment setup, running the notebook, and reproducing the main findings.
- Additional files detailing environment specifications (e.g., `requirements.txt` or `environment.yml`).

Clear documentation, organized notebook structure, and pinned package versions help ensure that others can replicate and extend our experiments effortlessly.

IV-D Initial Implementation Challenges

During early development, several implementation issues required iterative debugging and refinement:

¹<https://github.com/maverick250/research-project/blob/main/>

IV-D1 Stick-Breaking Construction

A misapplication of the stick-breaking weights initially caused $\sum_k \pi_k < 1$, violating the categorical distribution constraints. To clarify and resolve this issue, we revisit the relevant equations from Section III-B1a.

Recall that the stick-breaking process defines mixture weights as:

$$\pi_k = \beta_k \prod_{j=1}^{k-1} (1 - \beta_j), \quad k = 1, \dots, K_{\max}, \quad (19)$$

with $\beta_k \sim \text{Beta}(1, \alpha)$. Due to truncation at $K_{\max}$, the sum of weights did not reach one:

$$\sum_{k=1}^{K_{\max}} \pi_k < 1.$$

To rectify this, we add the remaining stick length to the last weight:

$$\pi_{K_{\max}} = \pi_{K_{\max}} + \prod_{k=1}^{K_{\max}} (1 - \beta_k). \quad (20)$$

This adjustment ensures that

$$\sum_{k=1}^{K_{\max}} \pi_k = 1.$$

Consequently, cluster assignments are sampled using the corrected mixture weights:

$$c_n \sim \text{Categorical}(\boldsymbol{\pi}_{\text{corrected}}), \quad n = 1, 2, \dots, N. \quad (21)$$

This correction prevents invalid log probabilities during sampling, improves convergence rates, and enhances the reliability of cluster assignments and latent feature allocations in the model.

IV-D2 Computational Scaling

Though feasible for moderate N and T , the MCMC sampling became increasingly time-consuming for larger simulated datasets or extended truncation levels. Possible mitigations (e.g., partial parallelization, advanced sampling algorithms) remain as future improvements for scaling to very large problems.

a) Impact on Model Reliability

Resolving these challenges permitted stable MCMC convergence, higher effective sample sizes, and consistent posterior inferences for cluster assignments and feature allocations. Consequently, the final implementation achieves both conceptual rigor (Section III) and robust practical performance on synthetic tests.

IV-E Summary of Implementation

By combining Python’s mature scientific ecosystem (PyMC, NumPy, Arviz, etc.) with a modular code design, we created a reproducible pipeline for data generation, model specification, and inference under a unified Bayesian nonparametric framework. This Implementation section ensures that future researchers can seamlessly re-run or extend our experiments, validate the findings, and adapt the approach to broader time-series applications.

V Results and Evaluation

This section presents the outcomes of the synthetic data analysis and evaluates the performance of the proposed Bayesian nonparametric model. The analysis covers the synthetic data overview, model inference results, model diagnostics, evaluation of clustering performance, feature allocation accuracy, parameter recovery, posterior predictive checks, generative capabilities, and computational performance.

V-A Synthetic Data Overview

A synthetic dataset comprising 50 time series was generated to evaluate the effectiveness of the proposed model. Each time series was assigned to one of three clusters and allocated latent features according to the Indian Buffet Process (IBP). The time series were simulated using a state-space model with cluster-specific parameters.

Figure 1 illustrates sample time series from each cluster, highlighting the distinct temporal dynamics corresponding to their cluster assignments.

The true cluster assignments and feature allocations serve as a reference for assessing the model’s inference capabilities.




Fig. 1: Sample Time Series from Each Cluster. Each subplot represents time series from a specific cluster, highlighting unique temporal behaviors.




Fig. 2: Original Heatmap of True Feature Allocation Matrix Z_{true} . Dark cells indicate active features across time series.




Fig. 3: True Feature Allocation Matrix Z_{true} Sorted by Cluster Assignments. Rows represent time series grouped by cluster, and columns represent features. Dark cells indicate feature activations.

V-A1 Feature Allocation Matrix Z_{true}

To visualize how latent features are distributed among time series and clusters, we present the true feature allocation matrix Z_{true} . Figure 2 displays the original heatmap of Z_{true} , while Figure 3 shows Z_{true} sorted based on cluster assignments for better interpretation. From Figure 3, the following observations regarding feature distribution across clusters can be made:

Observations:

- **Cluster 1:**
 - **Entirely Inactive Features:** Features 5, 12, and 16 show no activation (0) across all time series in Cluster 1.
 - **Predominantly Active Features:** Features 1, 3, 6, 7, 9, 11 (fully active), 13, 18, 19, and 20 exhibit high activation rates, with occasional deactivations.
- **Cluster 2:**
 - **Low Activation Features:** Within Cluster 2, Feature 9 exhibits the lowest activation, followed by relatively low activations in Features 2, 14, and 17.
 - **Predominantly Active Features:** Aside from these, most other features are highly activated, defining the characteristic latent feature profile of Cluster 2.
- **Cluster 3:**
 - **Low Activation Features:** Features 3, 7, and 11 have notably lower activation counts compared to others in Cluster 3.
 - **Predominantly Active Features:** Features 2, 4, 5, 10, 12 (fully active), 14, 15, 17 (fully active), and 20 (also fully active) are all highly activated, suggesting a heterogeneous latent structure.

Overall, the Z_{true} matrix confirms that the Indian Buffet Process (IBP) successfully allocated meaningful latent features, contributing to distinct temporal dynamics observed within each cluster.

Building on these observations of the true latent feature distributions, we now present the inference results, starting with the estimation of model parameters, followed by the assessment of feature and clustering allocations, and finally, an evaluation of the model’s generative capabilities.

VI Model Inference Parameters

This section examines the key parameters of the proposed Bayesian nonparametric model, focusing on their posterior estimates, convergence diagnostics (trace plots, $\hat{R}$ values, and effective sample sizes), and alignment with the true values used to generate the synthetic dataset. By assessing these parameters first, we establish whether the model has adequately recovered the core dynamics and whether additional sampling or refinement may be required.

VI-A Trace Plots and Posterior Summaries

Figure 4 displays the trace plots for a subset of key parameters, including the concentration parameters α_{dp} and α_p (for the Dirichlet Process and Indian Buffet Process respectively), the observation coefficient C , the process noise variance Q , the observation noise variance R , and the hyperparameter σ_{A_p} .




Fig. 4: Trace plots of selected parameters (α_{dp} , α_p , C , Q , R , σ_{A_p}). Each subplot shows the sampled values over MCMC iterations.

Visual inspection suggests that some parameters (notably α_p components and R) exhibit relatively stable mixing, whereas C and Q show signs of slower convergence, indicated by their chain trajectories and higher $\hat{R}$ values (as detailed below). Extended sampling or more informative priors may be necessary to improve mixing for these parameters.

VI-A1 Posterior Statistics

Table II provides a summary of the posterior distributions for α_{dp}, α_p , and the key parameters $\{C, Q, R, \sigma_{A_p}\}$. The $\hat{R}$ column indicates potential convergence issues for parameters with values significantly above 1.01–1.1, while the effective sample size (ESS) provides a measure of how many independent draws the MCMC sampler effectively obtained.

TABLE II: Posterior Means, Standard Deviations, 3% and 97% HDI Bounds, $\hat{R}$, and ESS for All Listed Parameters

Parameter	Mean	SD	3% HDI	97% HDI	ESS	$\hat{R}$
α_{dp}	1.157	0.757	0.165	2.554	4.0	1.42
$\alpha_p[0]$	3.029	1.440	0.876	5.837	73.0	1.01
$\alpha_p[1]$	1.662	1.365	0.058	4.257	15.0	1.10
$\alpha_p[2]$	1.752	2.515	0.010	6.619	3.0	1.89
$\alpha_p[3]$	2.226	1.471	0.261	4.992	68.0	1.03
$\alpha_p[4]$	1.883	1.358	0.117	4.349	24.0	1.07
$\alpha_p[5]$	1.783	1.476	0.083	4.436	8.0	1.17
$\alpha_p[6]$	1.869	1.327	0.176	4.412	127.0	1.01
$\alpha_p[7]$	1.618	1.561	0.032	4.385	4.0	1.44
$\alpha_p[8]$	2.103	1.398	0.244	4.658	195.0	1.01
$\alpha_p[9]$	1.838	1.424	0.116	4.514	38.0	1.05
C	0.379	0.112	0.227	0.608	3.0	1.93
Q	3.233	0.821	1.804	4.499	3.0	2.00
R	3.590	0.038	3.521	3.660	336.0	1.02
σ_{A_p}	0.028	0.024	0.000	0.074	4.0	1.54

a) Discussion of $\hat{R}$ and ESS

- **C** and **Q** exhibit $\hat{R}$ values of 1.93 and 2.00 respectively, suggesting that the MCMC chains have not fully converged or are mixing slowly for these parameters.
- Several of the α_p components (e.g., $\alpha_p[2]$, $\alpha_p[7]$) also exhibit elevated $\hat{R}$ (> 1.4) or low ESS, indicating potential challenges in exploring the IBP concentration parameter space.
- **R** stands out with $\hat{R} \approx 1.02$ and a large ESS (336.0), indicating more robust convergence and sampling efficiency for the observation noise variance.

VI-B Key Parameters vs. True Values

To further assess recovery of critical parameters, Figure 5 compares the posterior distributions of C , Q , R , and σ_{A_p} against their true values ($C_{\text{true}} = 2$, $Q_{\text{true}} = 0.5$, $R_{\text{true}} = 0.5$, $\sigma_{A_p, \text{true}} = 0.05$).




Fig. 5: Posterior distributions of key parameters (C, Q, R, σ_{A_p}) with true values indicated by vertical lines. Note the large discrepancy between true and estimated values for **C**, **Q**, and **R**.

VI-B1 Discussion of Parameter Recovery

- **Parameter Shifts:** The posterior means for C (≈ 0.38), Q (≈ 3.23), and R (≈ 3.59) deviate significantly from the true values (2.0, 0.5, 0.5). This suggests that the model is compensating for mismatches elsewhere - possibly in the latent feature allocations or cluster assignments - by inflating process and observation noise terms (i.e., Q and R) and lowering C .
- **Convergence Issues:** High $\hat{R}$ values for Q and C reinforce that the Markov chains may not have reached a stable posterior mode, thus the current posterior mean might not fully represent the model’s “final” estimate. Extending the number of samples, adjusting step sizes, or introducing more informative priors could help the chains stabilize.

- **Impact on the State-Space Model:** Overestimating Q implies that the model attributes more variability to the underlying state process, whereas an inflated R suggests higher observational noise. Both can blur distinctions between time series and degrade cluster resolution - potentially explaining misallocations in feature usage, discussed in subsequent sections.

In short, the model’s difficulty in recovering the true values indicates (1) potential under-specification or mismatch between the generative assumptions and actual data structure, and/or (2) insufficient MCMC exploration. Nevertheless, parameters such as $\mathbf{R}$ have converged well (high ESS, $\hat{R} \approx 1.02$) but at a value that greatly exceeds the true 0.5, revealing a systematic bias in how the noise is being modeled.

VI-C Implications for Further Analysis

Despite the divergence from true values for some parameters, these results still offer important insights:

- Tighter or more informative priors for Q and C could help constrain their ranges, preventing the sampler from drifting into large values that compensate for other model misfits.
- Additional MCMC iterations or a different sampling approach (e.g., variational inference) might be necessary to achieve better mixing.
- Since these parameters interact closely with the model’s latent structure, overshooting Q and R can dilute meaningful distinctions among time series. This, in turn, can adversely affect feature allocation and cluster assignments.

In the following sections, we will examine the model’s inferred feature allocations and clustering outcomes, where these parameter estimates play a pivotal role in shaping how the model partitions and explains the time series data.

VII Feature Allocation Results

Having examined the model’s key parameters and their convergence properties, we now turn to the inferred feature allocations. In our framework, each time series is associated with a subset of latent features determined via the Indian Buffet Process (IBP). These features then influence the time series’ underlying state dynamics through cluster-specific effect vectors. This section explores how well the inferred feature allocations ($\mathbf{Z}_{\text{inferred}}$) match the ground-truth allocations ($\mathbf{Z}_{\text{true}}$), the patterns observed in the sorted matrices, the degree of alignment measured by Jaccard Similarity scores, the cluster-specific feature activation probabilities, and the cluster-specific feature effects that govern state transitions.

VII-A True vs. Inferred Feature Allocation Matrices

Before comparing sorted matrices, we examine the original inferred feature allocation matrix.




Fig. 6: Original inferred feature allocation matrix $\mathbf{Z}_{\text{inferred}}$ before sorting. Dark cells indicate active features ($Z_{nf} = 1$).

Figure 6 shows the unsorted $\mathbf{Z}_{\text{inferred}}$ matrix directly from the model, reflecting the raw pattern of feature activations across all time series without any ordering by cluster.

Next, we sort both true and inferred matrices by cluster assignments to facilitate detailed comparison.

Figure 7 presents a comparison of the sorted true and inferred allocation matrices. Several notable discrepancies emerge:




Fig. 7: (Left) True feature allocation matrix $\mathbf{Z}_{\text{true}}$ sorted by cluster assignments. (Right) Inferred feature allocation matrix $\mathbf{Z}_{\text{inferred}}$ sorted in the same manner. Dark or colored cells indicate active features ($Z_{nf} = 1$).

- **Cluster 1 (Rows 1–23):** The $\mathbf{Z}_{\text{inferred}}$ matrix indicates universal activation of features (i.e., all 1s) for many time series in this cluster. While the ground truth also shows high feature activity, the model sometimes departs from subtle patterns (e.g., certain features that should be off for specific time series).
- **Cluster 2 (Rows 24–30):** Similar to Cluster 1, the inferred matrix indicates predominantly active features. However, a few time series (e.g., TS28, TS30) show some inactive features. Although this partially aligns with the true allocations, the overall pattern still leans toward “mostly active,” potentially missing finer granularity.
- **Cluster 3 (Rows 31–50):** The inferred matrix shows a nearly complete absence of features for all time series in this cluster ($Z_{nf} = 0$ for all f). This stark all-off pattern diverges from the ground truth whenever some features were actually active, implying the model has collapsed these time series into a “fully inactive” category.

a) Range of Estimated Feature Probabilities.

During inference, the model obtains a posterior probability $\hat{p}_{nf} \in [0.1875, 0.78]$ that each feature f is active in time series n . A threshold of 0.5 was then applied to convert these probabilities to binary allocations. This thresholding - combined with strong cluster-level effects - appears to have reinforced the all-on vs. all-off extremes.

VII-B Jaccard Similarity Analysis

To quantify the degree of overlap between true and inferred features on a per-time-series basis, we compute the Jaccard Similarity (JS):

$$\text{JS}(\mathbf{z}_{\text{true}}, \mathbf{z}_{\text{inferred}}) = \frac{|\mathbf{z}_{\text{true}} \cap \mathbf{z}_{\text{inferred}}|}{|\mathbf{z}_{\text{true}} \cup \mathbf{z}_{\text{inferred}}|}.$$

Selected Jaccard Similarities reveal that while some time series exhibit fairly high overlap (e.g., $\text{JS} \geq 0.80$), others show zero or near-zero similarity. Table III highlights a few examples:

TABLE III: Excerpt of Jaccard Similarity Scores for Selected Time Series

Time Series	Cluster	Jaccard Similarity
TS1	1	0.65
TS2	1	0.00
TS8	1	0.80
TS12	1	0.95
TS28	2	0.90
TS30	2	0.85
TS31	3	0.63
TS33	3	0.00
TS40	3	0.80
TS45	3	0.00
⋮	⋮	⋮
Mean		0.47
Median		0.68

Interpretation:

- **High Scores (≥ 0.80):** Some time series in Clusters 1 and 2 match their true features quite closely, indicating that when the model latches onto the correct on/off pattern, it can be highly accurate.
- **Zero Scores:** A substantial subset of time series (especially in Clusters 1 and 3) produce $JS=0$, reflecting no overlap in active features. Often, $\mathbf{Z}_{\text{true}}$ contained a specific partial pattern, but $\mathbf{Z}_{\text{inferred}}$ assigned either all-on or all-off.
- **Overall Mean (0.47) and Median (0.68):** The distribution is quite bimodal: some time series are well recovered, while many others are poorly matched. This misalignment may partially explain why process and observation noise parameters (Q and R) are inflated - compensating for incorrect feature usage.

VII-C Inferred Cluster-Specific Feature Activation Probabilities

Beyond the binary allocations, our model also infers the posterior mean activation probability $\mathbf{v}_{\text{clusters}}(k, f)$ for each cluster k and each feature f . Figure 8 plots these probabilities across features for up to 10 possible clusters (although in practice, we focus on the first 3 clusters, as all the time series were assigned to these clusters during inference).




Fig. 8: Inferred Cluster-Specific Feature Probabilities $\mathbf{v}_{\text{clusters}}(k, f)$. Each line represents one cluster, showing the average probability that feature f is active.

An excerpt of the numerical values for Clusters 1–3 is shown in Table IV (below), indicating:

- **Cluster 1** has high probabilities (≈ 0.71 to 0.76) for most features, consistent with near-complete activation in the inferred $\mathbf{Z}$.
- **Cluster 2** exhibits moderate probabilities (≈ 0.48 to 0.58), suggesting a more balanced on/off usage. However, thresholding at 0.5 often pushes many features to “on.”
- **Cluster 3** hovers around 0.35–0.38 for most features - slightly below the 0.5 threshold - hence the complete “off” pattern after thresholding.

TABLE IV: Illustrative Feature Activation Probabilities $\mathbf{v}_{\text{clusters}}(k, f)$ for Clusters 1–3 (Excerpt)

	Feat. 1	Feat. 2	Feat. 3	...	Feat. 19	Feat. 20
Cluster 1	0.7189	0.7289	0.7333	...	0.7463	0.7514
Cluster 2	0.5226	0.4867	0.5636	...	0.5465	0.5064
Cluster 3	0.3765	0.3701	0.3693	...	0.3738	0.3560

a) Discussion.

- **Strong Threshold Effect:** Because 0.5 is used as a universal cutoff, clusters with probabilities near 0.7+ end up fully activating most features, while clusters around 0.35 become almost entirely inactive.
- **Potential Soft Assignments:** In reality, a time series with probability 0.45 for a given feature may differ from one at 0.05 or 0.49, yet the binary “off” label lumps them together. Adjusting this threshold or using a soft assignment approach could capture more nuanced variation.

VII-D Inferred Cluster-Specific Feature Effects

Each cluster k has an associated feature-effect vector $\lambda_k \in \mathbb{R}^F$. When feature f is activated for a time series in cluster k , the autoregressive term A_n is shifted by $\lambda_{k,f}$. Figure 9 (placeholder) shows the posterior mean of these effects for each inferred cluster.




Fig. 9: Heatmap of Inferred Cluster-Specific Feature Effects, $\lambda_{A_clusters}(k, f)$. Positive (lighter) values amplify the state evolution, whereas negative (darker) values dampen it.

a) Notable Patterns:

- **Cluster 1:** Primarily small-to-moderate *positive* effects across most features ($\approx 0.13 \dots 0.19$), suggesting a gentle amplification of autoregressive dynamics in that cluster.
- **Cluster 2:** Predominantly *negative* feature effects (e.g., $-0.04 \dots -0.15$), implying a damping effect on time series transitions. This could partially explain differences in how time series evolve in Clusters 1 vs. 2.
- **Cluster 3:** Mostly modest or near-zero effects, consistent with the low probabilities (and near-zero allocations) of these features in the cluster. Because few features are “on,” the actual influence on the state equation is minimal.
- **Additional Clusters (4–10):** Even though some appear in the Dirichlet mixture, they often have smaller mixture weights. Their feature effects vary slightly (positive or negative) but have minimal impact on overall model behavior if few time series are assigned there.

VII-E Discussion of Feature Allocation and Effects

VII-E1 Interpretation of Feature Patterns

a) Over-Activation vs. Zero-Activation.

The model demonstrates a tendency to push entire clusters toward high or low probabilities, resulting in near-complete activation or deactivation of features after thresholding. While this strongly partitions time series into distinct clusters, it oversimplifies the “partial” usage of features that might exist in the true data.

b) Consequences for State-Space Dynamics.

Because each active feature shifts the autoregressive coefficient A_n by $\lambda_{c_n,f}$, misallocation (all-on or all-off) can lead to inflated noise terms (e.g., large Q, R) to explain unmodeled variability. Time series that truly required only a few selective features may be forced into a fully on/off scenario, biasing parameter estimates.

VII-E2 Alignment with Generative Assumptions

The observed patterns suggest that while the IBP prior successfully identifies a few dominant usage regimes, it struggles to represent moderate or irregular feature activity within each cluster. True data often exhibit more nuanced on/off behavior, requiring:

- **More informative priors on $v_{clusters}$ or λ_k :** This could discourage the model from converging to uniform extremes.
- **Adaptive Thresholding or Soft Assignments:** Instead of a hard 0.5 cutoff, the model or post-processing could treat probabilities near 0.4–0.6 differently, capturing partial or uncertain feature activation.

VII-E3 Implications for Clustering and Parameter Estimation

These feature-allocation outcomes are closely interwoven with the clustering mechanism (to be discussed in Section VIII). By oversimplifying feature usage, the model separates certain time series strongly (leading to high or zero Jaccard similarity) while inflating noise parameters to handle unexplained variability. Consequently, refining the fidelity of feature allocation is essential for improving time-series clustering and for aligning parameter estimates in the state-space model with their true values.

In the next section, we investigate how these feature patterns feed into the clustering process, the associated mixture weights, and ultimately, the model’s ability to partition the data coherently.

VIII Clustering Results

In this section, we evaluate the model’s ability to cluster the time series into distinct groups. We begin by examining the inferred mixture weights, which indicate the model’s posterior belief regarding how many clusters are actually “in use,” and then compare the true vs. inferred cluster assignments. Metrics such as the Adjusted Rand Index (ARI) are used to quantify clustering accuracy, and visual plots of time series grouped by inferred clusters provide further qualitative insights.

VIII-A Inferred Mixture Weights

The Dirichlet Process Mixture Model (DPMM) component of our approach produces a set of mixture weights, $\pi_{\text{corrected}}$, for up to a user-specified truncation level (here, 10 clusters). Figure 10 and Table V summarize these weights, revealing which clusters dominate.




Fig. 10: Inferred mixture weights $\pi_{\text{corrected}}$ for up to 10 clusters. The first three clusters hold the majority of the weight, with smaller fractions allocated to clusters 4–10.

TABLE V: Posterior Mean of Inferred Mixture Weights ($\pi_{\text{corrected}}$)

	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
π_k	0.3649	0.1537	0.2734	0.0613	0.0865	0.0335	0.0102	0.0058	0.0042	0.0066

a) Interpretation.

- **Dominance of First Three Clusters:** Clusters 1, 2, and 3 together capture $\approx 79\%$ of the total weight. This aligns with the known (true) scenario of three major clusters, although some smaller weights remain for clusters 4–10.
- **Residual Clusters:** Clusters 4–10 have nonzero weights but are collectively below 0.2. They likely serve as “spillover” or additional partitions that the DPMM introduces when the data suggest minor outliers or small variations.

While the high-level mixture weights confirm a strong focus on three main clusters, the next subsections detail whether the time series were actually assigned in a manner consistent with the true labels.

VIII-B True vs. Inferred Cluster Assignments

To assess the model’s clustering accuracy, we compare each time series’ true label to its most probable inferred label. Figure 11 provides an overview of these assignments. An excerpt of the assignments is shown in Table VI.

TABLE VI: Excerpt of True vs. Inferred Cluster Assignments

Index	True Cluster	Inferred Cluster
0	2	0
1	0	2
2	2	0
3	2	0
4	0	2
⋮	⋮	⋮
49	1	0




Fig. 11: Visualization comparing true and inferred cluster labels for each time series. Colors or markers can indicate the cluster ID assigned by the model vs. the ground truth.

a) Adjusted Rand Index.

The Adjusted Rand Index (ARI) measures the similarity between two clusterings while correcting for chance and is invariant to arbitrary label permutations. ARI values range from -1 to 1, where 1 indicates perfect agreement between the true and inferred clusterings, 0 implies random labeling, and negative values suggest less than chance agreement. This invariance means that the ARI does not depend on the specific numeric labels assigned to clusters, accommodating the arbitrary nature of label assignments in models like the DPMM.

We computed the ARI to assess the model’s clustering accuracy, obtaining a final **ARI** = 0.35, which suggests moderate-to-low agreement. This outcome can be attributed to the following causes:

- **Misclassification Patterns:** As seen in Figure 11, many time series that truly belong to cluster 2 or 3 are assigned to cluster 0 (i.e., the model’s first cluster). Meanwhile, some belonging to cluster 0 end up inferred as cluster 2, etc.
- **Influence of Feature Allocations:** Recall from Section VII that the model assigned near-complete on/off features in clusters 1 and 3, potentially blurring the differences among some time series. This likely contributes to the moderate ARI by binding together time series that share all-on (cluster 0) or all-off (cluster 2) patterns, regardless of subtle differences in the true data.

VIII-C Time Series Grouped by Inferred Clusters

Figure 12 shows representative time series from each inferred cluster. While the model creates broad groupings that partially reflect actual structures, discrepancies remain where feature allocations or high noise levels merge distinct true clusters.



Fig. 12: Example time series from each inferred cluster. Some clusters display coherent patterns, while others mix multiple true clusters.

a) *Observations.*

- **Coarse Separation:** Most time series that share broad amplitude or correlation traits end up together, but smaller differences are missed.
- **Influence of All-On or All-Off Features:** If clusters adopt near-complete activation/deactivation, the model can bind dissimilar series that share only these extremes in feature usage.

VIII-D Discussion of Clustering Performance

VIII-D1 Over- or Under-Clustering

Though the model can infer up to 10 clusters, it typically collapses around three. Some time series from a single true cluster are split across two inferred clusters or merged with other true clusters, indicating under-clustering on a finer scale.

VIII-D2 Model Limitations and Refinements

The moderate ARI suggests:

- **Better Feature Allocation:** Enhanced priors or thresholds could avoid the all-on/off pitfalls, improving cluster fidelity.
- **Noise Regularization:** Excessive variance in Q or R may camouflage important distinctions, hampering accurate assignments.
- **Extended Sampling:** More MCMC iterations or more advanced inference algorithms could yield more stable posterior modes, refining cluster separations.

VIII-E Summary of Clustering Outcomes

In summary, the DPMM approach identifies roughly three main clusters, aligning with the dominant data structures but failing to distinguish subtle subgroups. An ARI of 0.35 underscores the need for improved feature modeling, regularized noise parameters, or longer sampling to fully separate the true clusters. The following section examines posterior predictive checks and generative capabilities, offering further context on the model’s performance in synthesizing realistic time series.

IX Posterior Predictive Checks and Generative Capabilities

This final section assesses how well the proposed model reproduces the observed data (*posterior predictive checks*) and demonstrates its ability to generate new time series (*generative capabilities*). Both aspects provide insights into the model’s adequacy and limitations, especially given the challenges observed in parameter convergence and cluster/feature assignments.

IX-A Posterior Predictive Checks (PPC)

Posterior predictive checks compare actual observations to data simulated from the model’s posterior. We use `Arviz` to sample from the posterior predictive distribution and visualize the results.

IX-A1 Global PPC Plot

Figure 13 illustrates a global PPC for all time series. Although much of the observed data lie within the posterior predictive intervals, there is a noticeable discrepancy for certain time points. The observed data sometimes exceed the model’s central predictive range, especially in time series belonging to inferred Cluster 3, which exhibit more extreme oscillations or noise.

IX-A2 Representative PPC by Cluster

To pinpoint cluster-specific performance, we selected one representative time series from each inferred cluster and plotted its observed trajectory against the posterior predictive mean and ± 1 standard deviation. Figure 14 shows examples for three representative time series ($C = 1, 2, 3$).

a) *Observations.*

- **Cluster 1 & 2:** The observed trajectories generally fall within the model’s predicted bands, yielding moderate MAE and RMSE values (e.g., MAE ≈ 1.06 to 1.20 , RMSE ≈ 1.33 to 1.47).
- **Cluster 3:** The model struggles to capture the larger variance and oscillations, often underestimating or oversimplifying the series’ amplitude. Consequently, we see a much higher MAE (5.36) and RMSE (6.18) for the selected representative, indicating a mismatch between the model’s assumptions and the data’s volatility.

IX-A3 Quantitative Metrics and Coverage

`Arviz` summaries show stable sampling for most predictive draws ($\hat{R} \approx 1.0$), but the model’s overall coverage probability is ≈ 0.87 (below a nominal 94% HDI). This discrepancy again highlights difficulties in representing high-variance time series. For Clusters 1 and 2, MAE ranges from 1.06–1.20; in Cluster 3, it spikes to 5.36. RMSE follows a similar pattern, rising from ≈ 1.33 –1.47 to 6.18 for Cluster 3.

IX-B Generative Capabilities

Beyond fitting existing data, the model should also generate new time series that reflect each cluster’s characteristic patterns. We use the inferred cluster assignments, feature effect vectors $\lambda_{A_{\text{clusters}}}$, and cluster-specific feature activation probabilities v_{clusters} to simulate unseen time series.




Fig. 13: Posterior Predictive Check (Global). The distribution of predicted values is shown overlaid on the observed data for the entire dataset.

IX-B1 Synthesizing New Time Series by Cluster

Figure 15 presents several synthetic time series for each inferred cluster. The general shape, level of variance, and amplitude differ across clusters, roughly consistent with the model's parameter estimates:

- **Cluster 1:** The generated series exhibit moderate variance and relatively stable dynamics, though they do not perfectly mirror the original data's distinctive traits (likely due to misclassification issues and inflated noise).
- **Cluster 2:** These series show slightly higher oscillations or damping, in line with more negative feature effects (Section VII-D).
- **Cluster 3:** Despite an attempt to represent high variability, the synthetic series for Cluster 3 do not fully replicate the extreme oscillations in the observed data. The amplitude is less pronounced, suggesting that the model's inability to capture these signals in training similarly affects generation.

IX-B2 Discussion of Generative Performance

a) Reflecting Cluster-Specific Dynamics.

Overall, the generative samples align qualitatively with each cluster's broad features - damping or amplification of the autoregressive states - yet they are far from perfect replicas. The partial mismatch can be traced back to:

- **Overly Simplified Feature Use.** Clusters pushing all-on or all-off feature allocations, or clustering misassignments, ultimately shape the autoregressive coefficients incorrectly for some time series.
- **Inflated Noise Parameters.** As seen earlier, Q and R were substantially higher than their true values, diluting the distinctive patterns and leading to broader, less precise synthetic trajectories.

b) Implications for Real-World Application.

Should this method be applied to real data:

- **More Extensive Data or Stronger Priors** may help the model better separate genuine cluster distinctions from noise.




Fig. 14: Posterior Predictive Checks for one representative time series from each inferred cluster. Blue lines show observed data; red lines show posterior predictive means; gray shading represents ± 1 standard deviation.

- **Regularizing Noise Parameters** could prevent the model from “dumping” unexplained variance into Q and R .
- **Enhanced Feature Allocation Fidelity** (e.g., hierarchical priors or softer thresholds) could yield more nuanced generative behaviors, especially for clusters with moderate feature activations.

X Overall Summary of Results and Discussion

This concluding subsection synthesizes the findings across parameters, features, clustering, and generative checks:

a) Parameter Inference.

While some parameters (e.g., R) converged with high ESS and moderate $\hat{R}$, others (especially Q and C) showed slow mixing and diverged from their true values. The model may be compensating for misallocated features by inflating noise terms and shifting observation coefficients.




Fig. 15: Newly generated time series for Clusters 1, 2, and 3. Each subplot displays multiple generated realizations (5 per cluster).

b) Feature Allocation.

The IBP-driven mechanism tended to push clusters toward “all-on” or “all-off” patterns. Jaccard Similarities ranged from near 0.0 to 0.95 for different time series, leading to a bimodal distribution of accuracy.

c) Clustering Accuracy.

With an ARI of 0.35, clustering performance was moderate at best. Many time series from distinct true clusters were grouped together when they shared these oversimplified feature patterns, harming cluster recovery.

d) Posterior Predictive Checks.

PPC plots indicated that while Clusters 1 and 2 were reasonably well-represented (MAE 1.06–1.20, RMSE 1.33–1.47), Cluster 3 remained problematic (MAE 5.36, RMSE 6.18). Coverage probability reached only 0.87, below the targeted HDI.

e) Generative Capabilities.

Newly synthesized series reflected some broad cluster differences in variance and dynamics but failed to capture the more extreme oscillations of Cluster 3 or the finer nuances of Cluster 1.

f) Methodological Lessons and Future Steps.

- More iterative or robust inference (longer MCMC or variational approaches) may help reduce parameter drift and improve feature/cluster alignments.
- A refined thresholding strategy or hierarchical priors on $\mathbf{v}_{\text{clusters}}$ could capture partial feature usage more accurately, avoiding the all-on/all-off extremes.
- Regularizing or constraining noise parameters (Q, R) could prevent excessive inflation, encouraging the model to explain variability through latent features instead of pure noise.

Overall, the proposed Bayesian nonparametric state-space framework demonstrates the potential for flexible clustering and feature allocation in time-series data, but its performance on this synthetic test reveals a few systematic tendencies - most notably, oversimplified feature usage and inflated noise parameters - that warrant further refinement before application to complex real-world datasets.

XI Limitations

Despite the contributions presented in this research, several limitations should be acknowledged:

- **Data Scale and Complexity:** The synthetic dataset used in this study does not fully capture the richness and variability of real-world time series. Although synthetic data allow controlled experiments, real datasets often exhibit irregular sampling, seasonality, complex noise structures, and abrupt regime changes, which our controlled setup does not replicate. As a result, model performance on genuine, large-scale data may deviate significantly and must be evaluated further.
- **Inference Challenges:** Certain parameters, particularly the process noise variance Q and the observation coefficient C , displayed high $\hat{R}$ values, indicating slow MCMC mixing or potential multi-modality in the posterior. This suggests that extended sampling, more robust priors, or alternative inference approaches (e.g., variational inference) might be needed for stable convergence. Additionally, the high noise estimates may obscure cluster and feature distinctions, impacting model interpretability.
- **Feature Over-Simplification:** The model’s thresholding procedure within the Indian Buffet Process (IBP) led to near-complete on/off feature activations in some clusters. While this cleanly separated certain groups, it likely overlooks nuanced partial feature usage. Over-simplified feature allocations can inflate noise parameters (Q and R) as the model compensates for unexplained variability, and may reduce the fidelity of inferred cluster distinctions.
- **Computational Burden:** MCMC-based inference is computationally intensive. As the number of time series, features, or time steps grows, scaling to real-world applications can become challenging. Further optimization - such as parallelization or advanced sampling schemes - may be required for feasibility in large-scale scenarios.

These limitations underscore the need for cautious interpretation of the results and provide a roadmap for potential methodological improvements in future research.

XII Future Work

Building on the findings of this research, multiple avenues exist for enhancing the proposed Bayesian nonparametric state-space model:

- **Refined Prior Structures:** Introducing hierarchical or structured priors on the cluster-specific feature probabilities $\mathbf{v}_{\text{clusters}}$ could curb the tendency toward extreme (all-on/all-off) allocations. Correlated priors or sparsity-inducing priors might also help capture more realistic feature usage patterns.
- **Variational Inference or Extended MCMC:** Investigating alternative inference strategies - such as Variational Inference (VI), Gibbs sampling with more sophisticated block updates, or parallel-tempering approaches - may address slow mixing for critical parameters like Q and C . These methods might also reduce computational overhead relative to standard Hamiltonian Monte Carlo or NUTS.

- **Real-World Validation:** Applying the framework to large-scale real-world datasets in domains such as finance, healthcare, or signal processing would test its robustness under more diverse conditions (e.g., missing data, structural breaks, nonstationary behavior). Such validation would inform hyperparameter choices, refine prior specifications, and guide further model extensions.
- **Regularized Noise Parameters:** Imposing stronger priors or hierarchical structures on noise variances (Q and R) could prevent the model from “dumping” unmodeled variability into noise terms. Regularization may encourage the model to utilize latent features more effectively, potentially improving both clustering accuracy and feature allocation.
- **Temporal Extensions and Switching Dynamics:** Future research could explore variants such as sticky HDP-HMMs or dynamic IBP formulations that allow cluster and feature usage to evolve over time. These approaches might better capture regime shifts or time-varying behaviors in complex real-world time series.

Implementing these enhancements could lead to more accurate parameter estimation, richer feature allocations, and improved clustering performance, ultimately broadening the applicability of the model to various data-intensive fields.

XIII Conclusion

In this research, we presented a Bayesian nonparametric state-space framework aimed at clustering and feature allocation in multiple time series without specifying the number of clusters or features *a priori*. Our approach integrated the Dirichlet Process Mixture Model (DPMM) to enable flexible and unbounded clustering, and the Indian Buffet Process (IBP) to allocate an unbounded set of latent features. These Bayesian nonparametric tools were embedded in a state-space model to capture both the global partitioning of time series into clusters and the local, feature-driven dynamics within each cluster.

Key Contributions and Findings

- **Bayesian Nonparametric Integration:** We unified the DPMM and IBP within a state-space context, providing a principled means of simultaneously learning cluster assignments and latent features. This framework accommodates scenarios where the true number of clusters or features is unknown or expected to grow with the data.
- **Synthetic Data Experiments:** Using controlled synthetic datasets, we demonstrated that the model can uncover broad cluster-level structures and can generate synthetic time series that loosely mirror the underlying cluster dynamics. However, the experiments revealed a tendency toward misallocated features, often resulting in overestimated noise parameters (Q and R).
- **Clustering and Feature Allocation Performance:** Although the model achieved some success in separating time series into coherent groups, the Adjusted Rand Index (ARI) of roughly 0.35 highlights the challenges stemming from uniform on/off allocations. The posterior predictive checks showed that Clusters 1 and 2 were relatively well-modeled, while Cluster 3 posed significant difficulties due to high variance and oscillations.
- **Model Diagnostics and Convergence:** Posterior summaries indicated that certain parameters (C and Q) did not converge as smoothly (high $\hat{R}$), suggesting multi-modal posteriors or insufficient sampling. These challenges underscore the complexities in nonparametric priors and the need for more sophisticated inference strategies or extended sampling runs.

Implications and Next Steps

The results confirm the promise of Bayesian nonparametric methods for flexible time-series clustering but also highlight the pitfalls of threshold-based IBP allocations and inflated noise terms. Future work could address these issues via (1) improved priors that promote partial feature usage rather than extreme allocations, (2) more extensive or alternative sampling methods to resolve posterior multi-modality, and (3) rigorous testing on larger, real-world datasets that contain the complexities absent in synthetic data.

In conclusion, this thesis/research work provides both a proof-of-concept and a cautionary tale. On the one hand, it demonstrates that a unified DPMM–IBP state-space approach can reveal clusters and features in time-series data with minimal structural assumptions. On the other hand, the observed mismatches, slow mixing, and feature oversimplifications imply that further refinements, particularly in inference and prior design, are essential. By building on the lessons learned here - refining feature allocation, regularizing noise, and exploring advanced sampling techniques - future research can harness the full potential of Bayesian nonparametric modeling in multi-series time-domain applications.

References

- [1] J. Bovy, D. Terenin, and B. Davidson, “ArviZ: Exploratory analysis of Bayesian models in Python,” *arXiv preprint arXiv:1809.09444*, 2018.
- [2] J. Durbin and S. J. Koopman, *Time Series Analysis by State Space Methods*. Oxford University Press, 2002.
- [3] M. D. Fox, Y. W. Teh, B. Kulis, and A. Welling, “Sticky hierarchical Dirichlet processes,” *Proceedings of the 26th Annual International Conference on Machine Learning (ICML)*, 2009, pp. 1441–1448.
- [4] C. Fraley and A. E. Raftery, “Model-based clustering, discriminant analysis, and density estimation,” *Journal of the American Statistical Association*, vol. 97, no. 458, pp. 611–631, 2002.
- [5] A. Gelman and D. B. Rubin, “Inference from iterative simulation using multiple sequences,” *Statistical Science*, vol. 7, no. 4, pp. 457–472, 1992.
- [6] Z. Ghahramani and M. E. Tipping, “Infinite mixtures of infinite Gaussian mixtures,” *Proceedings of the 20th International Conference on Neural Information Processing Systems (NIPS)*, 2005, pp. 1137–1144.

- [7] I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio, “Generative adversarial nets,” *Proceedings of the 27th International Conference on Neural Information Processing Systems (NIPS)*, 2014, pp. 2672–2680.
- [8] T. Griffiths and M. Ghahramani, “The infinite latent feature model and the Indian buffet process,” *Canadian Journal of Statistics*, vol. 39, no. 2, pp. 293–314, 2011.
- [9] C. R. Harris, K. Millman, R. van der Walt, J. Gommers, R. Virtanen, A. Cournapeau, B. Wieser, A. Taylor, J. Berg, R. Smith, K. Kern, M. Picus, R. Hoyer, J. van Kerkwijk, M. Brett, F. H. Haldane, and P. Cournapeau, “Array programming with NumPy,” *Nature*, vol. 585, pp. 357–362, 2020.
- [10] T. Hastie, R. Tibshirani, and J. Friedman, *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*, 2nd ed. Springer, 2009.
- [11] G. E. Hinton and R. Salakhutdinov, “Reducing the dimensionality of data with neural networks,” *Science*, vol. 313, no. 5786, pp. 504–507, 2006.
- [12] S. Hochreiter and J. Schmidhuber, “Long short-term memory,” *Neural Computation*, vol. 9, no. 8, pp. 1735–1780, 1997.
- [13] M. Hoffman, D. Blei, D. Wang, and B. Paisley, “No-U-Turn Sampler: Adaptively setting path lengths in Hamiltonian Monte Carlo,” *Journal of Machine Learning Research*, vol. 15, pp. 1593–1623, 2014.
- [14] L. A. Hubert and P. Arabie, “Comparing partitions,” *Journal of Classification*, vol. 2, no. 1, pp. 193–218, 1985.
- [15] J. D. Hunter, “Matplotlib: A 2D graphics environment,” *Computing in Science & Engineering*, vol. 9, no. 3, pp. 90–95, 2007.
- [16] R. J. Hyndman and G. Athanasopoulos, *Forecasting: principles and practice*, 2nd ed. OTexts, 2018.
- [17] R. J. Hyndman and G. Athanasopoulos, *Forecasting: Principles and Practice*, 2nd ed. OTexts, 2018.
- [18] H. Ishwaran and E. James, “Stick-breaking representations of Dirichlet processes,” *Journal of the American Statistical Association*, vol. 96, no. 453, pp. 161–173, 2001.
- [19] R. Kumar, O. Martin, and A. L. Rodriguez, “ArviZ: A Python library for exploratory analysis of Bayesian models,” *Journal of Open Source Software*, vol. 4, no. 33, p. 1143, 2019.
- [20] V. M. Mäkeläinen, “Advanced Bayesian Nonparametric Methods for Time Series Clustering,” *arXiv preprint arXiv:1803.11299*, 2018.
- [21] V. Malik and J. Jain, “A survey of clustering algorithms for time series data,” *IEEE Transactions on Knowledge and Data Engineering*, vol. 22, no. 1, pp. 144–158, 2010.
- [22] M. Mirza and S. Osindero, “Conditional generative adversarial nets,” *arXiv preprint arXiv:1411.1784*, 2014.
- [23] R. M. Neal, “Markov Chain Sampling Methods for Dirichlet Process Mixture Models,” *Journal of Computational and Graphical Statistics*, vol. 9, no. 2, pp. 249–265, 2000.
- [24] A. Nikolaev, M. Smith, and L. Johnson, “Synthetic data generation for time series analysis: techniques and applications,” *Journal of Synthetic Data*, vol. 5, no. 2, pp. 123–145, 2021.
- [25] OpenAI, “ChatGPT: Language Model for Conversations,” 2023. [Online]. Available: <https://chat.openai.com>. [Accessed: January 11, 2025].
- [26] The PyTorch Development Team, “PyTorch: A tensor library for probabilistic programming,” Available: <https://pytorch.readthedocs.io>, 2023.
- [27] A. Radford, L. Metz, and S. Chintala, “Unsupervised representation learning with deep convolutional generative adversarial networks,” *arXiv preprint arXiv:1511.06434*, 2015.
- [28] J. Salvatier, A. Wiecki, and A. Fonnesbeck, “Probabilistic programming in Python using PyMC3,” *PeerJ Computer Science*, vol. 2, p. e55, 2016.
- [29] R. H. Shumway and D. S. Stoffer, *Time Series Analysis and Its Applications: With R Examples*, 4th ed. Springer, 2010.
- [30] A. Smith and B. Johnson, “TSGM: A framework for generative modeling of synthetic time series,” *arXiv preprint arXiv:2305.11567*, 2023.
- [31] Y. W. Teh, B. Kulis, C. Jordan, and A. Welling, “Hierarchical Dirichlet processes,” *Journal of the American Statistical Association*, vol. 101, no. 476, pp. 1566–1581, 2006.
- [32] Y. W. Teh, J. S. Blei, and M. I. Jordan, “A collapsed Gibbs sampler for latent Dirichlet allocation,” *Journal of Machine Learning Research*, vol. 11, pp. 581–618, 2010.
- [33] S. van Rossum and F. L. Drake Jr., *Python Tutorial: A Quick Introduction*, Available: <https://docs.python.org/3/tutorial/>, 2011.
- [34] T. Vinh, N. Epps, and J. Bailey, “Information-theoretic measures for clusterings comparison: Variants, properties, normalization and correction for chance,” *arXiv preprint arXiv:1002.2663*, 2010.
- [35] Pierre, “Bayesian Nonparametric Model for Time Series Clustering,” GitHub repository, Available: <https://github.com/maverick250/research-project/blob/main/>, 2024.